

Cache-Oblivious Search Trees via Trees of Small Height

Gerth Stølting Brodal

 BRICS

University of Aarhus

Joint work with Rolf Fagerberg and Riko Jacob

Result : New Search Tree

$\{ 1, 3, 4, 5, 6, 7, 8, 10, 11, 13 \}$

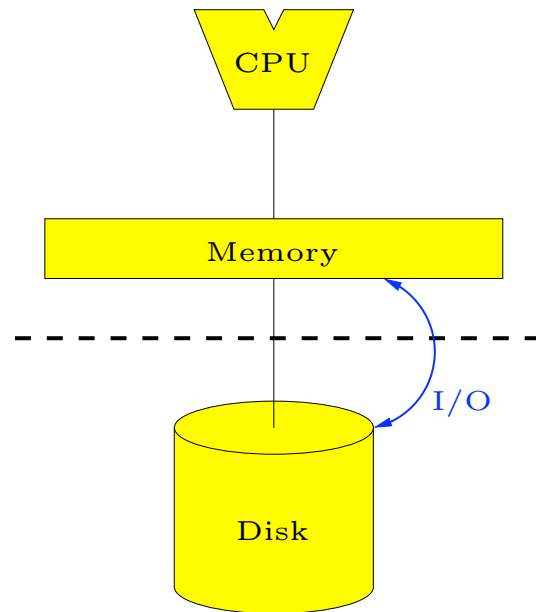


| | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|----|----|----|
| 6 | 4 | 8 | 1 | — | 3 | 5 | — | — | 7 | — | — | 11 | 10 | 13 |
|---|---|---|---|---|---|---|---|---|---|---|---|----|----|----|

Outline

- ▶ ● Models of Computation
 - I/O Model
 - Cache-Oblivious Model
- Cache-Oblivious Search Trees
 - Static
 - Dynamic
- Experiments
 - Memory Layouts of Trees
- Summary

I/O Model



N = problem size

M = memory size

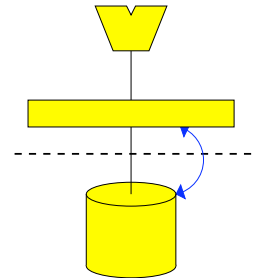
B = I/O block size

Aggarwal and Vitter 1988

- Bottleneck \equiv I/Os between the two highest memory levels
- B-trees support searches and updates in $O(\log_B N)$ I/Os
- $\Theta\left(\frac{M}{B}\right)$ -way merge-sort achieves optimal $\Theta\left(\frac{N}{B} \log_{M/B} \frac{N}{B}\right)$ I/Os

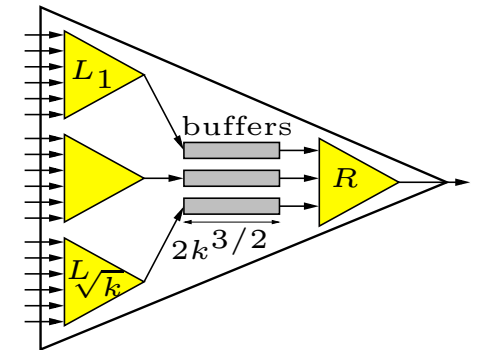
Cache-Oblivious Model

- I/O model
- Algorithms **do not know** the parameters B and M
- Optimal off-line cache replacement strategy



Examples

- Scanning, Linear time selection
- Matrix-transposition, FFT, Funnel-sorting



Lemma

Optimal cache-oblivious algorithm implies optimal algorithm on **each level** of a **fully associative** multi-level cache using LRU

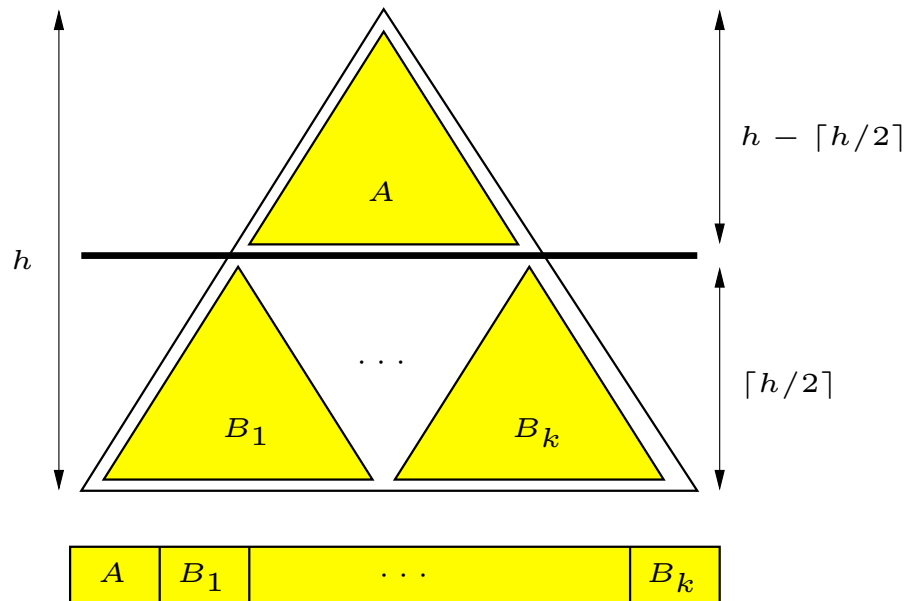
Frigo et al. 1999

Outline

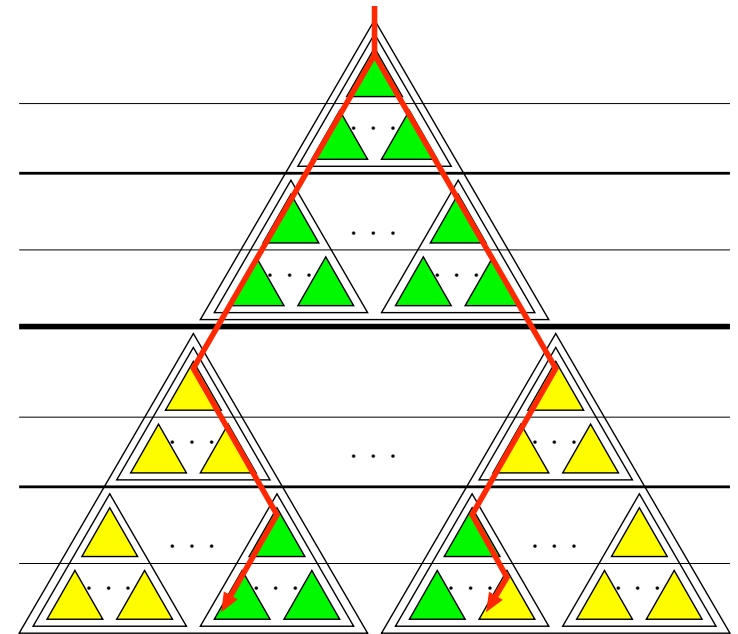
- Models of Computation
 - I/O Model
 - Cache-Oblivious Model
- ▶ • Cache-Oblivious Search Trees
 - Static
 - Dynamic
- Experiments
 - Memory Layouts of Trees
- Summary

Static Cache-Oblivious Trees

Recursive memory layout \equiv van Emde Boas layout



Degree $O(1)$



Searches use $O(\log_B N)$ I/Os

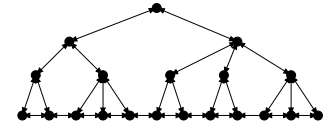
Range reportings use

$O\left(\log_B N + \frac{k}{B}\right)$ I/Os

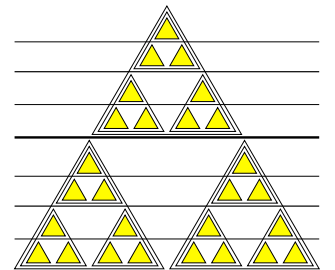
Prokop 1999

Dynamic Cache-Oblivious Trees

| | |
|-----------------|---|
| Search | $O(\log_B N)$ |
| Range Reporting | $O\left(\log_B N + \frac{k}{B}\right)$ |
| Updates | $O\left(\log_B N + \frac{\log^2 N}{B}\right)$ |



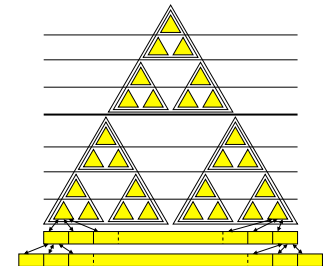
Arge and Vitter 1996



Prokop 1999



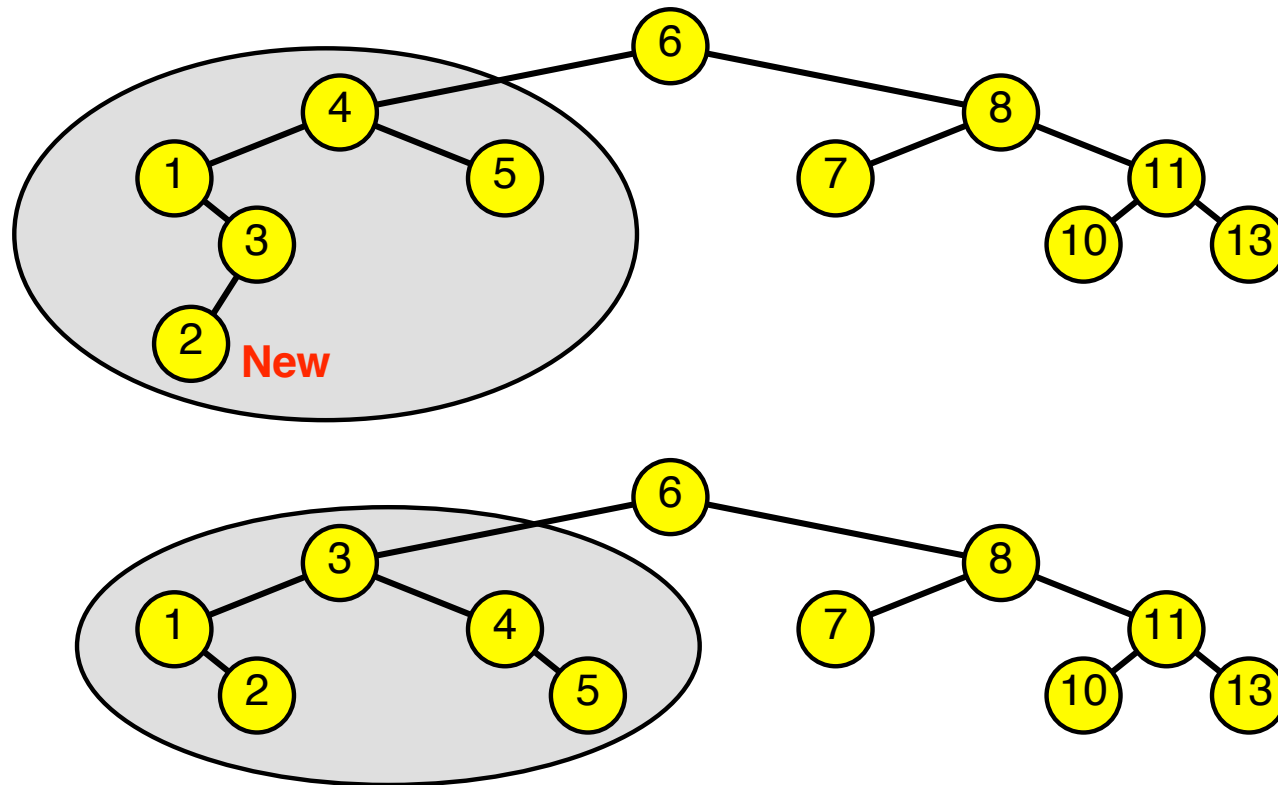
Itai et al. 1981



Bender, Demain, Farach-Colton 2000

- Pointer Based Strongly Weight Balanced B-trees
- Dynamic van Emde Boas Layout
- Packed Memory Management
- Two Levels of Indirection

Binary Trees of Small Height

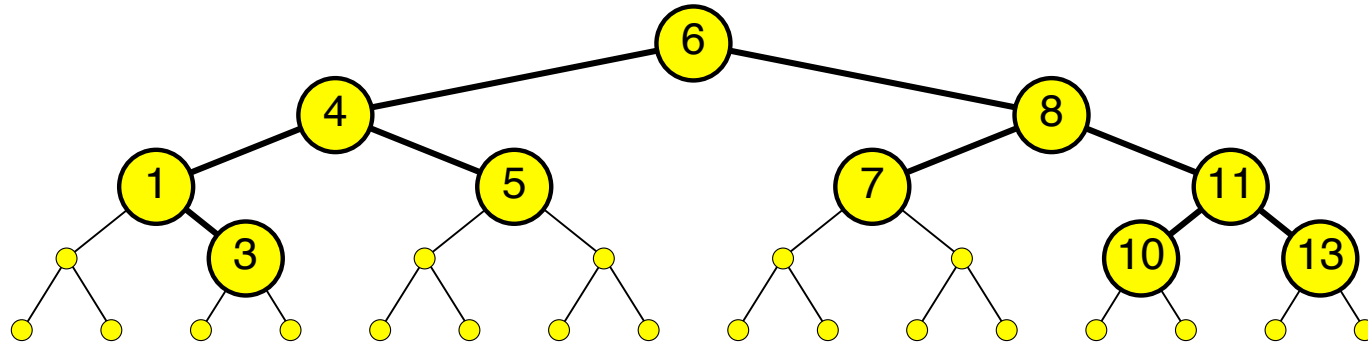


- If an insertion causes non-small height then **rebuild subtree** at nearest ancestor with sufficient few descendents
- Insertions require amortized time $O(\log^2 N)$

Andersson and Lai 1990

Dynamic Cache-Oblivious Trees

- **Embed** a dynamic tree of small height into a complete tree
- **Static** van Emde Boas layout



Search

$$O(\log_B N)$$

Range Reporting

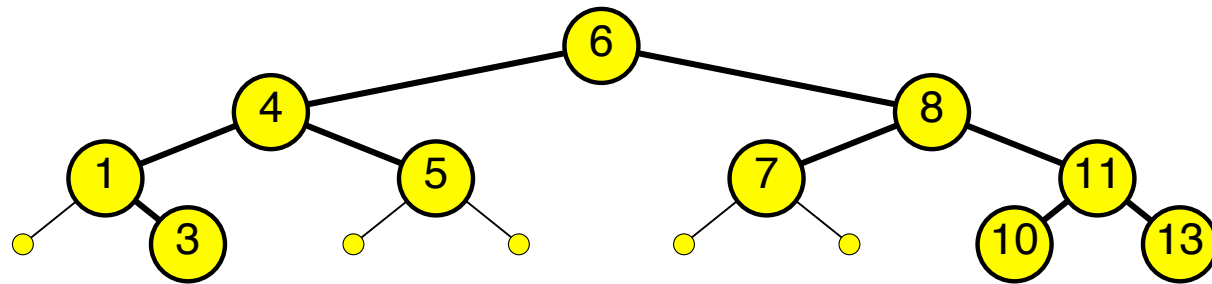
$$O\left(\log_B N + \frac{k}{B}\right)$$

Updates

$$O\left(\log_B N + \frac{\log^2 N}{B}\right)$$

New

Example

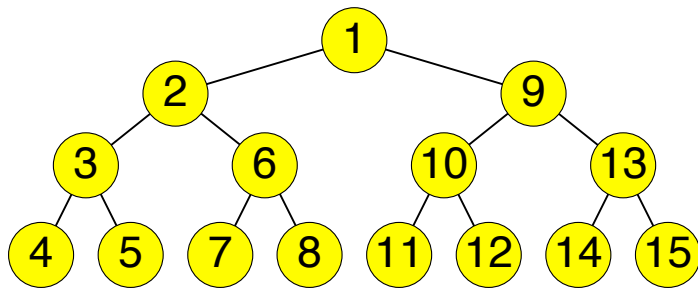


| | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|----|----|----|
| 6 | 4 | 8 | 1 | — | 3 | 5 | — | — | 7 | — | — | 11 | 10 | 13 |
|---|---|---|---|---|---|---|---|---|---|---|---|----|----|----|

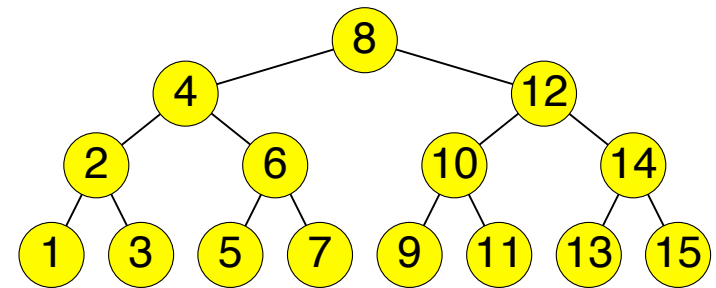
Outline

- Models of Computation
 - I/O Model
 - Cache-Oblivious Model
- Cache-Oblivious Search Trees
 - Static
 - Dynamic
- ▶ • Experiments
 - Memory Layouts of Trees
- Summary

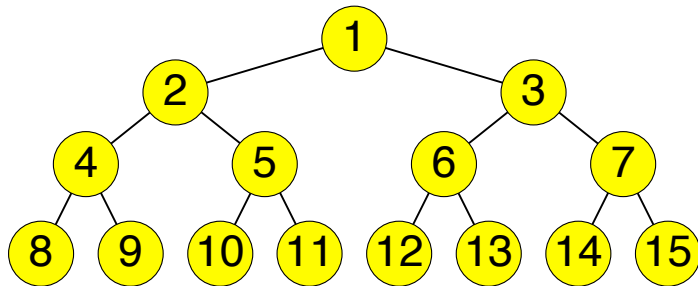
Memory Layouts of Trees



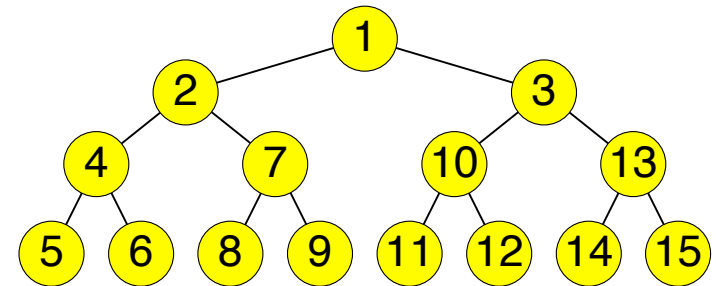
DFS



inorder

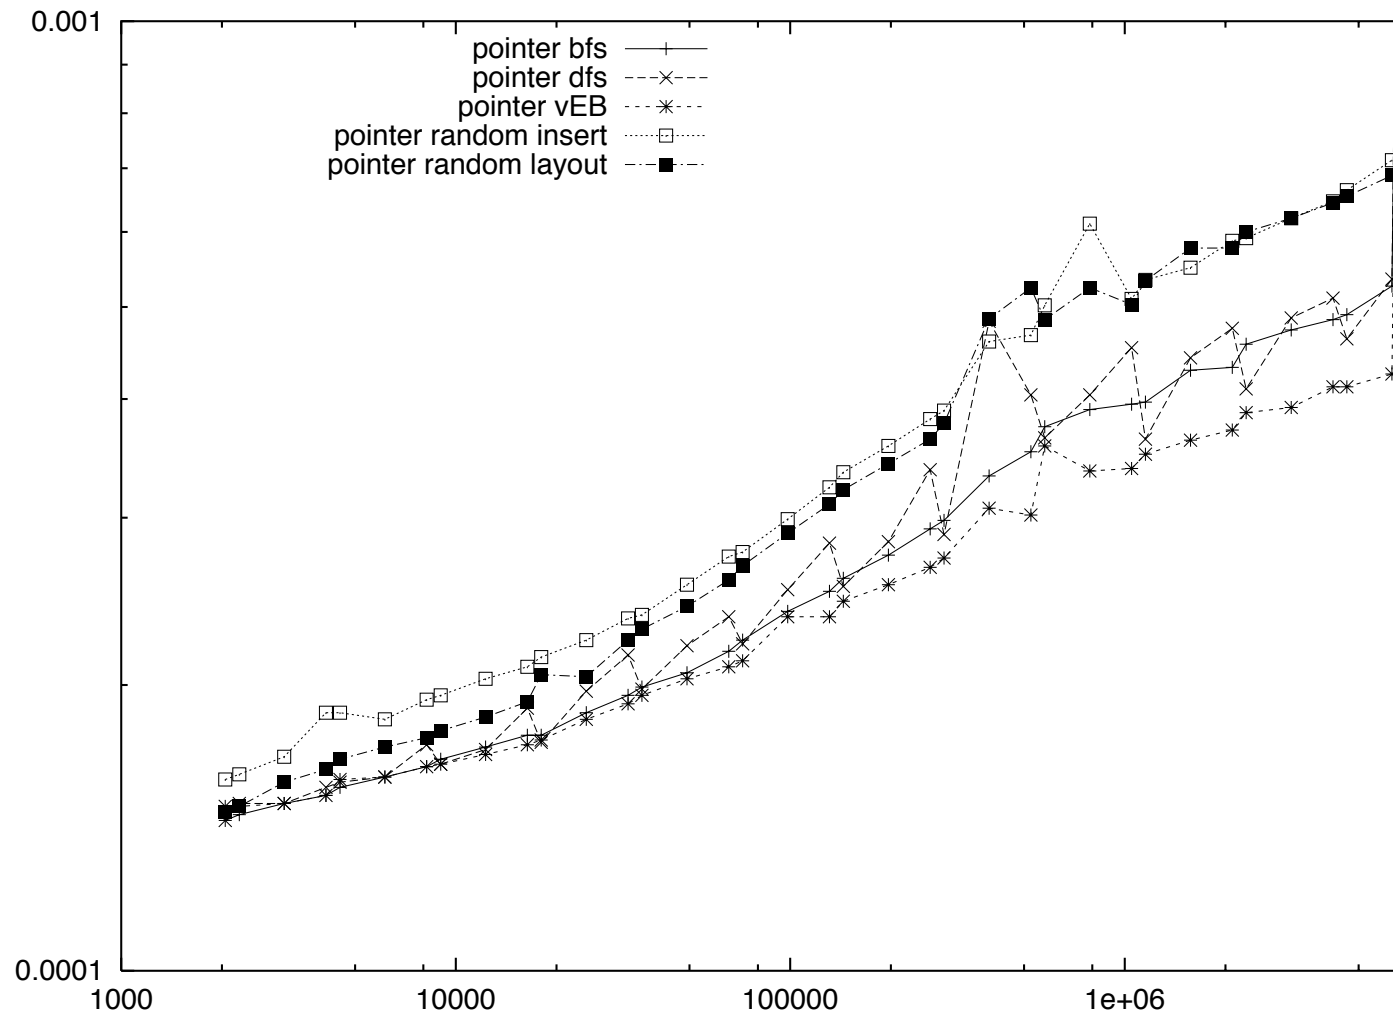


BFS



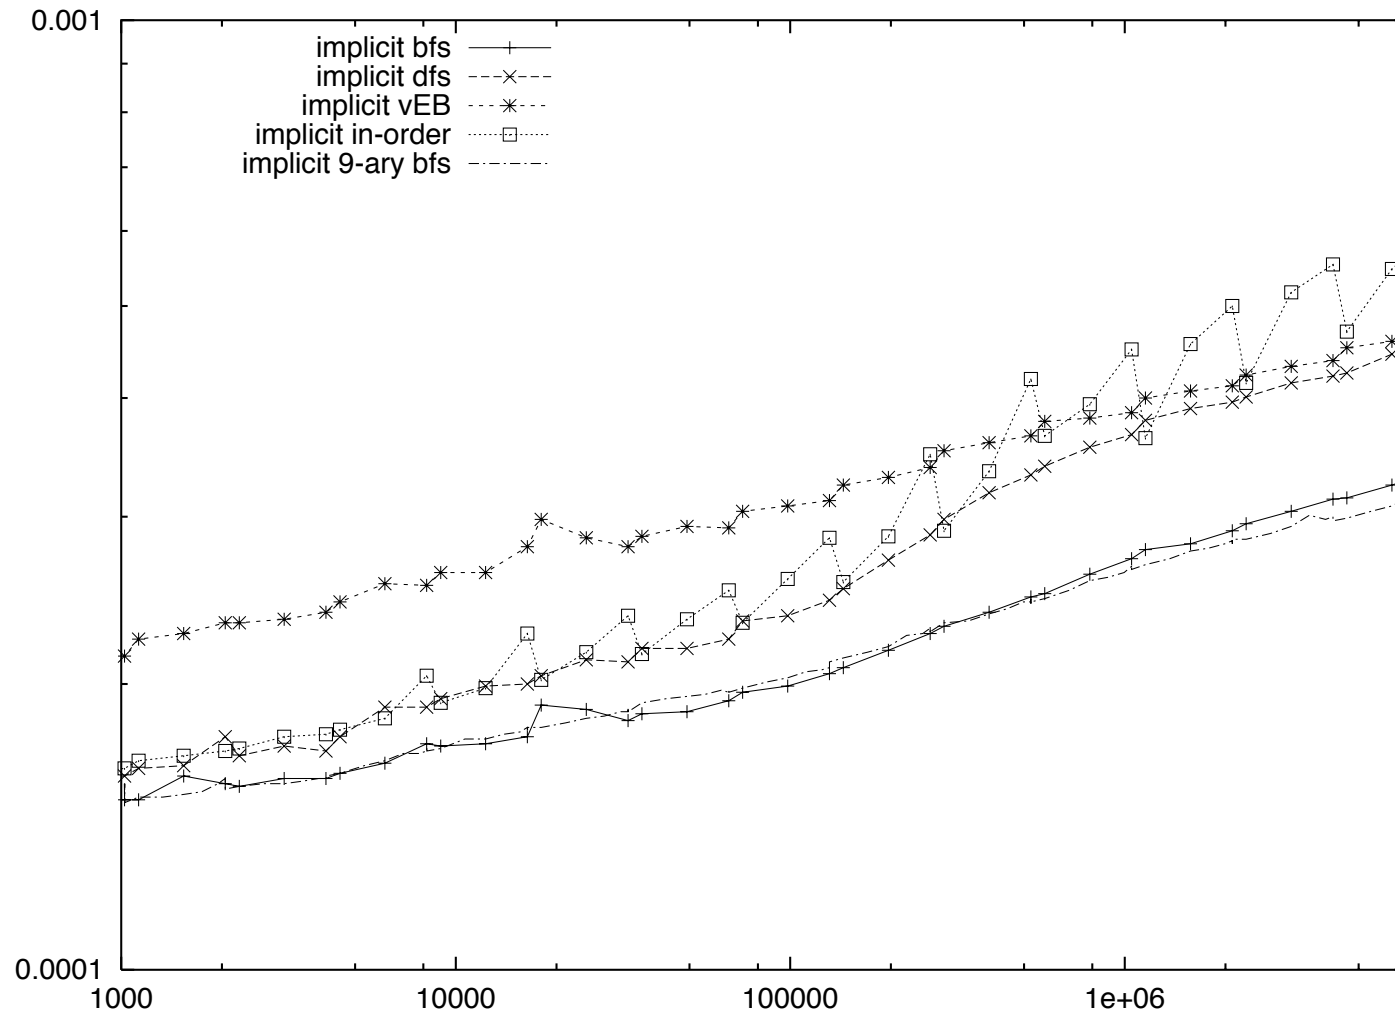
van Emde Boas

Searches in Pointer Based Layouts



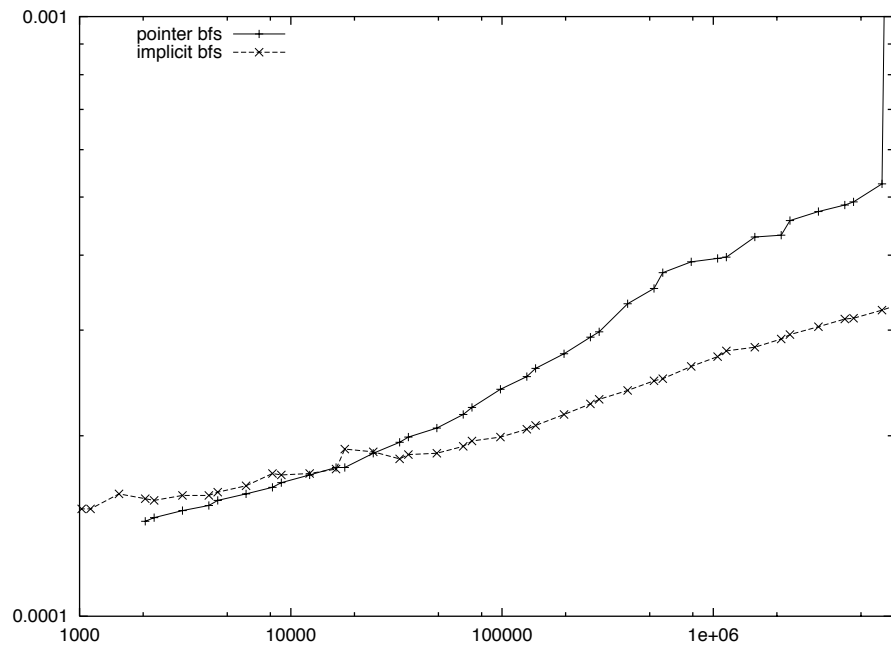
- van Emde Boas layout wins, followed by the BFS layout

Searches with Implicit Layouts

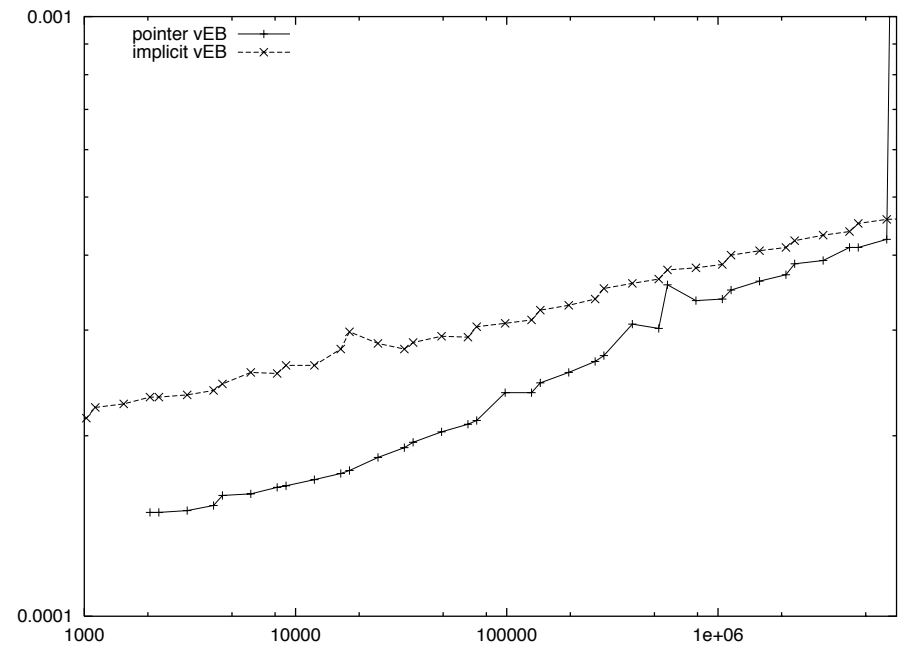


- BFS layout wins due to simplicity and caching of topmost levels
- van Emde Boas layout requires quite complex index computations

Implicit vs Pointer Based Layouts



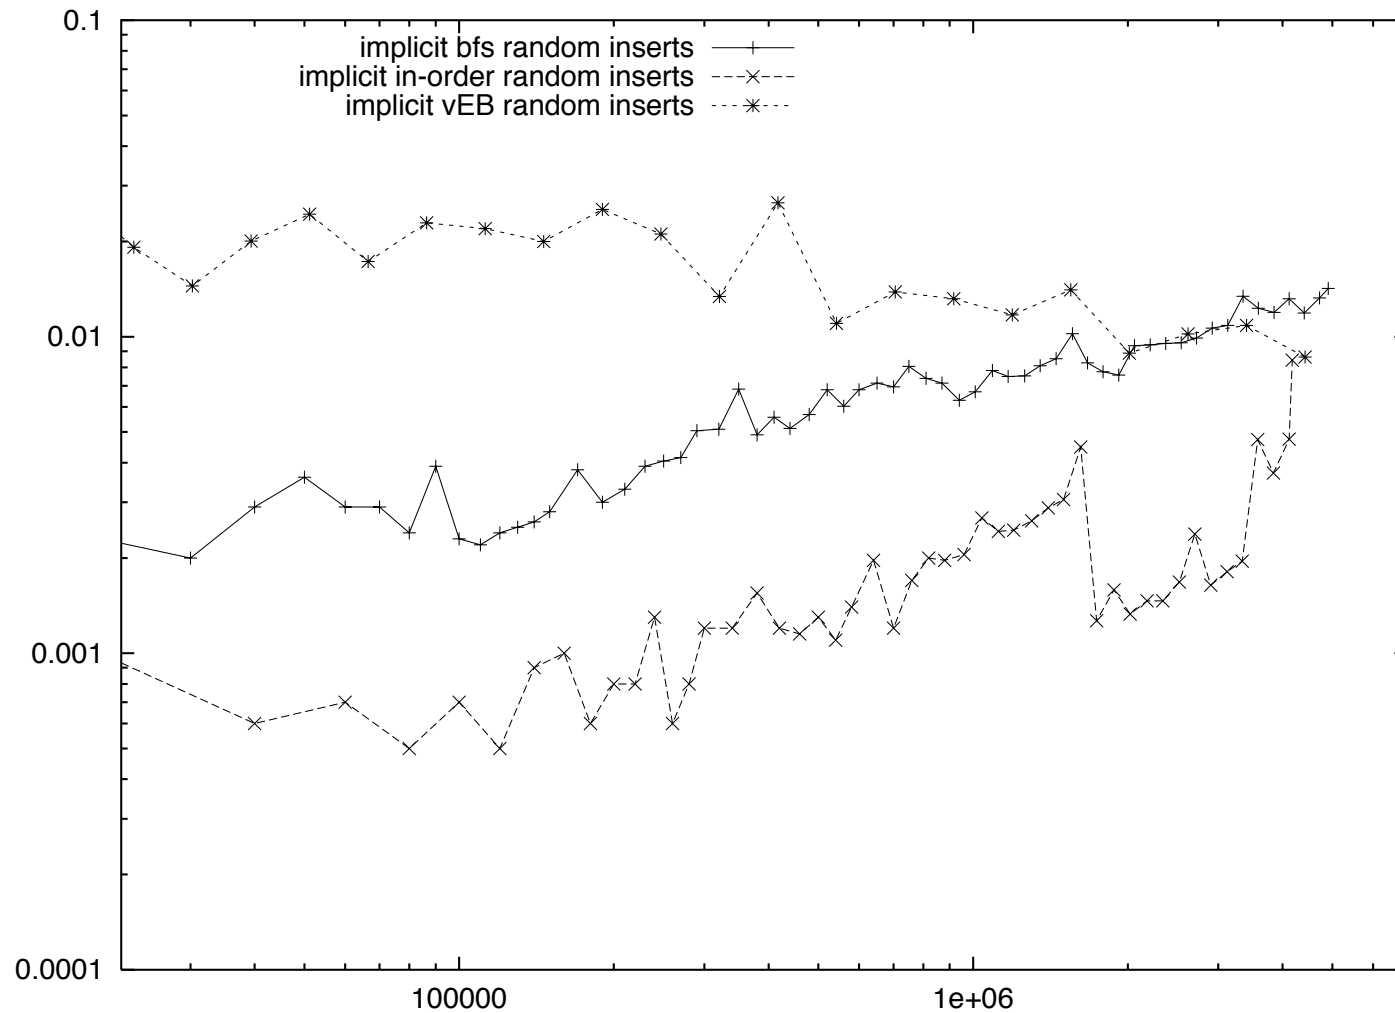
BFS layout



van Emde Boas layout

- Implicit layouts become competitive as n grows

Insertions in Implicit Layouts



- Insertions are rather slow (factor 10-100 over searches)

Summary

- New simple cache-oblivious search trees

| | |
|-----------------|---|
| Search | $O(\log_B N)$ |
| Range Reporting | $O\left(\log_B N + \frac{k}{B}\right)$ |
| Updates | $O\left(\log_B N + \frac{\log^2 N}{B}\right)$ |

- Update time $O(\log_B N)$ by one level of indirection (implies sub-optimal range reporting)
- Importance of memory layouts
- van Emde Boas layout gives good cache performance
- Computation time is important when considering caches

| | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|----|----|----|
| 6 | 4 | 8 | 1 | — | 3 | 5 | — | — | 7 | — | — | 11 | 10 | 13 |
|---|---|---|---|---|---|---|---|---|---|---|---|----|----|----|