

# Cache Oblivious Search Trees via Binary Trees of Small Height

Gerth Stolting Brodal  
Rolf Fagerberg  
Riko Jacob

# Lecture Outline

- Motivation
- General idea & Working methods
- van Emde Boas memory layout
- The algorithm
  - Dictionary operations
  - Insertion complexity proof
- Experiments & Results
- Summary

# Motivation

- Modern computers contain a hierarchy of memory levels.
- Access time is about 1 [cyc] for registers/L1 cache, and 100,000 [cyc] for the disk.
- The cost of a memory access depends highly on what is the lowest memory level containing the data.
- The evolution of CPU speed and memory access time indicates that these differences are likely to increase in the future.

# Motivation

- Our goal:
  - To find an implementation for **binary search tree** that tries to minimize cache misses.
  - That algorithm will be **cache oblivious**.
- By optimizing an algorithm to one **unknown** memory level, it is optimized to each memory level automatically !

# General idea & Working methods

- Definitions:
- A tree  $T_1$  can be **embedded** in another tree  $T_2$ , if  $T_1$  can be obtained from  $T_2$  by pruning subtrees.
- **Implicit layout** - the navigation between a node and its children is done based on address arithmetic, and not on pointers.

# General idea & Working methods

- Assume we have a binary search tree.
- Embed this tree in a **static complete** tree.
- Save this (complete) tree in the memory in a cache oblivious fashion.
- Whenever  $n$  is doubled we create a new static tree.

# General idea & Working methods

## ➤ Advantages:

- Minimizing memory transfers.
- Cache obliviousness
- No pointers – better space utilization:
  - A larger fraction of the structure can reside in lower levels of the memory.
  - More elements can fit in a cache line.

## ➤ Disadvantages:

- Implicit layout: higher instruction count per navigation – slower.


# van Emde Boas memory layout

- Recursive definition:
- A tree with only one node is a single node record.
- If a tree  $T$  has two or more nodes:
  - Divide  $T$  to a top tree  $T_0$  with height  $\lceil h(T)/2 \rceil$  and a collection of bottom trees  $T_1, \dots, T_k$  with height  $\lceil h(T)/2 \rceil$ , numbered from left to right.
  - The van Emde Boas layout of  $T$  consist of the v.E.B. layout of  $T_0$  followed by the v.E.B. layout of  $T_1, \dots, T_k$




# van Emde Boas memory layout

➤ Example :




# van Emde Boas memory layout

➤ Example :




# van Emde Boas memory layout

➤ Example :




# van Emde Boas memory layout

## ➤ Example :




# van Emde Boas memory layout

## ➤ Example :




# van Emde Boas memory layout

➤ Example :




# van Emde Boas memory layout

➤ Example :



# van Emde Boas memory layout


➤ Example :






# van Emde Boas memory layout

➤ Example :




# van Emde Boas memory layout

➤ Example :



# van Emde Boas memory layout

➤ Example :



# The algorithm

## ➤ Search:

- Standard search in a binary tree.
- Memory transfers:  $O(\log_B n)$  worst case

## ➤ Range query:

- Standard range query in a binary tree:
  - Search the smallest element in the range
  - Make an inorder traversals till you reach an element greater then or equals to the greatest element in the range.
- Memory transfers:  $O(\log_B n + k/B)$  worst case

# The algorithm

## ➤ Notations:

$T$  = the dynamic tree.

$H$  = the height of the static complete tree.

$s(v)$  = the size of the subtree in the complete tree rooted at  $v$ .

$\rho(v) = \frac{|T_v|}{s(v)}$  = the density of  $v$ .

# The algorithm

We'll define a sequence of evenly distributed density thresholds:  $0 < \tau_1 < \tau_2 < \dots < \tau_H = 1$

by :  $\tau_i = \tau_{i-1} + \Delta$

$$\Delta = \frac{(1 - \tau_1)}{(H - 1)}$$

Example:  $H = 5$

$$\tau_1 = 0.6$$

$$\tau_2 = 0.7$$

$$\tau_3 = 0.8$$

$$\tau_4 = 0.9$$

$$\tau_5 = 1$$

$$\Rightarrow \Delta = 0.1$$

# The algorithm

**Invariant:** for the root  $r$ :

$$\rho(r) \leq \tau_1 \Rightarrow H \geq \log\left(\frac{n}{\tau_1} + 1\right)$$

## ➤ Insertions:

- Locate the position in  $T$  of the new node (regular search)
- If  $d(v) = H+1$  we rebalance  $T$

# Rebalancing

1. Find the nearest ancestor  $w$  of  $v$  with:

$$\rho(w) < \tau_{d(w)}$$

- That means: find the nearest ancestor of  $v$  which is not too dense.
- In order to do so we have to calculate  $|T_w|$
- We can do it by a simple traversal – Why?



# Rebalancing


2. After having located  $w$  we rebalance  $w$  as follows:

- Create a sorted array of all elements in  $T_w$  by an inorder traversal of  $T_w$ .
- The middle element in the array stored at  $w$ .
- The smallest (greatest) half elements are recursively redistributed at the left (right) child of  $w$ .

# Rebalancing

➤ Example : insert 7


$$\begin{aligned}\tau_1 &= 0.8, \\ \tau_2 &= 0.9, \\ \tau_3 &= 1\end{aligned}$$



# Rebalancing

➤ Example : insert 7

$$\begin{aligned}\tau_1 &= 0.8, \\ \tau_2 &= 0.9, \\ \tau_3 &= 1\end{aligned}$$




# Rebalancing

➤ Example : insert 7

$$\begin{aligned}\tau_1 &= 0.8, \\ \tau_2 &= 0.9, \\ \tau_3 &= 1\end{aligned}$$

4	5	6	7	8	10
---	---	---	---	---	----


$$\rho = 0.71$$



# Rebalancing

➤ Example : insert 7


$$\begin{aligned}\tau_1 &= 0.8, \\ \tau_2 &= 0.9, \\ \tau_3 &= 1\end{aligned}$$



# Rebalancing

➤ Example : insert 7


$$\begin{aligned}\tau_1 &= 0.8, \\ \tau_2 &= 0.9, \\ \tau_3 &= 1\end{aligned}$$



# Rebalancing

➤ Example : insert 7

$$\begin{aligned}\tau_1 &= 0.8, \\ \tau_2 &= 0.9, \\ \tau_3 &= 1\end{aligned}$$




# Rebalancing

➤ Example : insert 7

$$\tau_1 = 0.8,$$

$$\tau_2 = 0.9,$$

$$\tau_3 = 1$$



➤ The next insertion will cause a rebuilding



# Insertion complexity

- **Lemma:** A redistribution at  $v$  implies

$$\lfloor \rho(v) \cdot s(w) - 1 \rfloor \leq |T_w| \leq \lceil \rho(v) \cdot s(w) \rceil$$

for all descendants  $w$  of  $v$

- In other words: after a redistribution at  $v$ ,  
for all descendant  $w$

$$\rho(w) \cong \rho(v)$$

- Proof: (induction)

# Insertion complexity

➤ Theorem:

Insertions require amortized  $O(\frac{\log^2 n}{1 - \tau_1})$  time

and amortized  $O(\log_B n + \frac{\log^2 n}{B(1 - \tau_1)})$  memory transfers

# Insertion complexity (time)

## ➤ Proof:

- Consider a distribution at a node  $v$ , caused by an insertion below  $v$ .
- $\Rightarrow$  for a child  $w$  of  $v$  :

$$\rho(w) \geq \tau_{d(w)} \Rightarrow |T_w| \geq \tau_{d(w)} \cdot s(w)$$

# Insertion complexity (time)

## ➤ Proof (cont.):

- The Lemma argues that immediately after a redistribution at  $v$ , for all descendant  $w$  of  $v$ :

$$|T_w| \leq \lceil \rho(v) \cdot s(w) \rceil$$

- Since the redistribution took place at  $v$ :

$$\rho(v) < \tau_{d(v)}$$

•  $\Rightarrow$

$$|T_w| \leq \tau_{d(v)} \cdot s(w) + 1$$

# Insertion complexity (time)

## ➤ Proof (cont.):

- It follows that the number of insertions below  $w$  since the last redistribution at  $v$  or an ancestor of  $v$  is at least:

$$\tau_{d(w)} \cdot s(w) - (\tau_{d(v)} \cdot s(w) + 1)$$

The number of elements in  $w$  right now, because  $w$  become “dense”

The number of elements at  $w$  immediately after the last redistribution at  $v$  or at ancestor of  $v$

# Insertion complexity (time)

➤ Proof (cont.):

$$\tau_{d(w)} \cdot s(w) - (\tau_{d(v)} \cdot s(w) + 1)$$

$$= \tau_{d(w)} \cdot s(w) - \tau_{d(v)} \cdot s(w) - 1$$

$$= s(w) (\tau_{d(w)} - \tau_{d(v)}) - 1$$

$$= s(w) \cdot \Delta - 1$$

# Insertion complexity (time)

## ➤ Proof (cont.):

- The redistribution at  $v$  takes  $O(s(v))$  which can be covered by  $\max\{s(v), \Lambda - 1\}$

Hence, each node is charged at:

$$O\left(\frac{s(v)}{\max\{1, s(w) \cdot \Lambda - 1\}}\right) = O\left(\frac{1}{\Lambda}\right)$$

for each of the mentioned insertions below  $w$ .

# Insertion complexity (time)

## ➤ Proof (cont.):

- Since each node has at most  $H$  ancestors it will be charged at most  $H$  times and the amortized complexity will be:

$$H \cdot O\left(\frac{1}{\Delta}\right) = O\left(\frac{H}{\Delta}\right) = O\left(\frac{H^2}{1 - \tau_1}\right) = O\left(\frac{\log^2 n}{1 - \tau_1}\right)$$



# Insertion complexity (memory transfers)

## ➤ Proof:

- A top down search requires  $O(\log_B n)$  memory transfers.
- The redistribution itself is like a range query when  $k = c(v)$ , hence takes:  
 $O(\log_B n + \frac{\log^2 n}{B(1-\tau_1)})$  memory transfers.
- Amortized got:

Finding the ancestor

Redistribution at the ancestor

$$O(\log_B n + \frac{\log^2 n}{B(1-\tau_1)})$$

# Experiments & Results

## ➤ The platform:


- Two 1[GHz] Pentium 3 processors.
- 256[KB] of cache
- 1[GB] R.A.M.
- Operating system: Linux kernel 2.4.3-12smp.

## ➤ The experiments – **searches** with:

- Various tree sizes (different  $n$ )
- Various memory layouts
- Implicit and pointer based versions


# Experiments & Results

- Red: cache aware algs.
- Blue: v.E.B layout.
- Green: BFS implicit
- Black: inorder implicit



# Experiments & Results

- Red: v.E.B pointer
- Green: v.E.B implicit
- Black: B.F.S. pointer
- Blue: B.F.S. implicit



# Summary

- We introduced the van Emde Boas layout.
- We learned a cache oblivious search tree implementation and its time & memory transfers bounds.
- We saw that this implementation is indeed efficient and competitive with cache aware implementations.

# Thank you all

14/6/04

Elad Levi